

INFORME TECNICO

Introduccion

El presente informe técnico se enfoca en analizar y comparar estrategias de optimización de consultas en SQL, centrándose en el uso de SQL básico, FAV (window functions), y la cláusula WITH (common table expressions). Se han planteado diversas problemáticas basadas en un caso de estudio (Biblioteca Nacional), evaluando los tiempos de ejecución de cada enfoque para determinar la opción más eficiente.

Desarrollo

FAV - WITH Benchmark

Se crearon 5 problemáticas del caso de estudio (Biblioteca Nacional) y de cada uno se mostrara una posible implementacion con sql basico, fav, y la clausula with, en conjunto con los tiempos medidos

1. Crear 5 problemáticas del caso de estudio
 1. aplicar sql basico, fav, with

2. explain analyse (screenshot)
3. escoger el mejor

ejemplo 1

Se necesita de los prestamos que se entregan en un dia darle prioridad a los que han estado mas tiempo prestados, para eso:

Se desea mostrar los datos de los prestamos (id_document, id_service, fecha inicio, fecha entrega), asi como la prioridad de los prestamos que se entregan en un mismo dia con 1 la mas alta que tendra el que mas tiempo a estado prestado (pueden haber empates)

SQL Basico

```
select id_document,  
       id_service,  
       start_date,  
       end_date,  
       (select count(distinct term) + 1  
        from loan l2  
        where l2.end_date = l1.end_date
```

```
and l2.term > l1.term) as priority  
from loan l1;
```

	❏ QUERY PLAN	⌵
1	Seq Scan on loan l1 (cost=0.00..1416599.38 rows=7944 width=24) (actual tim...	
2	SubPlan 1	
3	-> Aggregate (cost=178.29..178.31 rows=1 width=8) (actual time=0.294...	
4	-> Sort (cost=178.26..178.28 rows=7 width=4) (actual time=0.292...	
5	Sort Key: l2.term	
6	Sort Method: quicksort Memory: 25kB	
7	-> Seq Scan on loan l2 (cost=0.00..178.16 rows=7 width=4)...	
8	Filter: ((term > l1.term) AND (end_date = l1.end_date...	
9	Rows Removed by Filter: 7933	
10	Planning Time: 0.071 ms	
11	JIT:	
12	Functions: 13	
13	Options: Inlining true, Optimization true, Expressions true, Deforming tr...	
14	Timing: Generation 0.508 ms, Inlining 12.745 ms, Optimization 33.794 ms, ...	
15	Execution Time: 2408.406 ms	

FAV

```
select loan.id_document,  
       loan.id_service,  
       loan.start_date,
```

```

        loan.end_date,
        rank() over (partition by end_date order by term desc) as priority
from loan
order by end_date desc;

```

	❏ QUERY PLAN	⌵
1	Sort (cost=1326.52..1346.38 rows=7944 width=28) (actual time=9.082..9.453 ...	
2	Sort Key: end_date DESC	
3	Sort Method: quicksort Memory: 627kB	
4	-> WindowAgg (cost=653.04..811.92 rows=7944 width=28) (actual time=3.08...	
5	-> Sort (cost=653.04..672.90 rows=7944 width=20) (actual time=3.0...	
6	Sort Key: end_date, term DESC	
7	Sort Method: quicksort Memory: 565kB	
8	-> Seq Scan on loan (cost=0.00..138.44 rows=7944 width=20) ...	
9	Planning Time: 0.045 ms	
10	Execution Time: 10.050 ms	

WITH

```

explain analyze
with ranking as (select id_service, id_document,
                        rank() over (partition by end_date order by term desc) as
priority
                        from loan)

```

```
select loan.id_document, loan.id_service, start_date, end_date, priority
from loan
      join ranking on loan.id_service = ranking.id_service
      and loan.id_document = ranking.id_document;
```

	❏ QUERY PLAN	⌵
1	Hash Join (cost=910.64..1190.66 rows=1 width=24) (actual time=5.858...	
2	Hash Cond: ((loan_1.id_service = loan.id_service) AND (loan_1.id_doc...	
3	→ WindowAgg (cost=653.04..811.92 rows=7944 width=24) (actual time...	
4	→ Sort (cost=653.04..672.90 rows=7944 width=16) (actual tim...	
5	Sort Key: loan_1.end_date, loan_1.term DESC	
6	Sort Method: quicksort Memory: 503kB	
7	→ Seq Scan on loan loan_1 (cost=0.00..138.44 rows=794...	
8	→ Hash (cost=138.44..138.44 rows=7944 width=16) (actual time=2.54...	
9	Buckets: 8192 Batches: 1 Memory Usage: 437kB	
10	→ Seq Scan on loan (cost=0.00..138.44 rows=7944 width=16) (...)	
11	Planning Time: 0.225 ms	
12	Execution Time: 13.088 ms	

Resultados: FAV esa la mejor opcion por simplicidad y velocidad

ejemplo 2

Se quiere hacer un censo de los miembros que son estudiantes universitarios jóvenes asociados a la biblioteca así como el promedio de edad por país de los estudiantes

De ellos se desea obtener los datos (nombre, edad, país, escuela) así como el promedio de edad

SQL Basico

```
select m1.id_member,
       name,
       age,
       country,
       school,
       (select avg(age)
        from member m2
        where m1.country = m2.country
          and m1.category = 'student') as average_age
from member m1
      join student on m1.id_member = student.id_member
where age between 20 and 40
      and school like 'University%';
```

	❏ QUERY PLAN	⌵
1	Hash Join (cost=1059.97..13008552.64 rows=3529 width=99) (actual time=237...	
2	Hash Cond: (m1.id_member = student.id_member)	
3	-> Seq Scan on member m1 (cost=0.00..4109.15 rows=42858 width=39) (actu...	
4	Filter: ((age >= 20) AND (age <= 40))	
5	Rows Removed by Filter: 126988	
6	-> Hash (cost=884.50..884.50 rows=14038 width=36) (actual time=6.425..6...	
7	Buckets: 16384 Batches: 1 Memory Usage: 1064kB	
8	-> Seq Scan on student (cost=0.00..884.50 rows=14038 width=36) (a...	
9	Filter: ((school)::text ~~ 'University% '::text)	
10	Rows Removed by Filter: 28388	
11	SubPlan 1	
12	-> Aggregate (cost=3684.68..3684.69 rows=1 width=32) (actual time=17...	
13	-> Result (cost=0.00..3682.97 rows=685 width=4) (actual time=0...	
14	One-Time Filter: (m1.category = 'student '::category_type)	
15	-> Seq Scan on member m2 (cost=0.00..3682.97 rows=685 wid...	
16	Filter: ((m1.country)::text = (country)::text)	
17	Rows Removed by Filter: 127700	
18	Planning Time: 0.352 ms	
19	JIT:	
20	Functions: 26	
21	Options: Inlining true, Optimization true, Expressions true, Deforming tr...	
22	Timing: Generation 0.994 ms, Inlining 21.128 ms, Optimization 111.775 ms,...	
23	Execution Time: 63055.212 ms	

FAV

```
select member.id_member,  
       name,  
       age,  
       country,  
       school,  
       avg(age) over (partition by country) as average_age  
from member  
       join student on member.id_member = student.id_member  
where age between 20 and 40  
       and school like 'University%';
```


	❏ QUERY PLAN	⌵
1	WindowAgg (cost=5489.58..5551.34 rows=3529 width=99) (actual time=28.615...	
2	-> Sort (cost=5489.58..5498.40 rows=3529 width=67) (actual time=28.591...	
3	Sort Key: member.country	
4	Sort Method: quicksort Memory: 446kB	
5	-> Hash Join (cost=1059.97..5281.63 rows=3529 width=67) (actual t...	
6	Hash Cond: (member.id_member = student.id_member)	
7	-> Seq Scan on member (cost=0.00..4109.15 rows=42858 width=...	
8	Filter: ((age >= 20) AND (age <= 40))	
9	Rows Removed by Filter: 126988	
10	-> Hash (cost=884.50..884.50 rows=14038 width=36) (actual t...	
11	Buckets: 16384 Batches: 1 Memory Usage: 1064kB	
12	-> Seq Scan on student (cost=0.00..884.50 rows=14038 ...	
13	Filter: ((school)::text ~~ 'University% '::text)	
14	Rows Removed by Filter: 28388	
15	Planning Time: 0.275 ms	
16	Execution Time: 30.826 ms	

WITH

```
with average_age_per_country as (select country, avg(age) as average_age
                                from member
                                group by country)

select member.id_member,
       name,
       age,
```

```
        member.country,  
        school,  
        average_age  
from member  
        join student on member.id_member = student.id_member  
        join average_age_per_country on member.country =  
average_age_per_country.country  
where school like 'University%';
```

	❏ QUERY PLAN	⌵
1	Hash Join (cost=5167.19..8908.95 rows=14038 width=99) (actual time=41.680...	
2	Hash Cond: ((member.country)::text = (average_age_per_country.country)::...	
3	-> Hash Join (cost=1059.97..4764.26 rows=14038 width=67) (actual time=...	
4	Hash Cond: (member.id_member = student.id_member)	
5	-> Seq Scan on member (cost=0.00..3256.77 rows=170477 width=35) ...	
6	-> Hash (cost=884.50..884.50 rows=14038 width=36) (actual time=6...	
7	Buckets: 16384 Batches: 1 Memory Usage: 1064kB	
8	-> Seq Scan on student (cost=0.00..884.50 rows=14038 width...	
9	Filter: ((school)::text ~~ 'University% '::text)	
10	Rows Removed by Filter: 28388	
11	-> Hash (cost=4104.10..4104.10 rows=249 width=40) (actual time=35.432...	
12	Buckets: 1024 Batches: 1 Memory Usage: 22kB	
13	-> Subquery Scan on average_age_per_country (cost=4068.62..4104...	
14	-> Finalize GroupAggregate (cost=4068.62..4101.61 rows=249...	
15	Group Key: member_1.country	
16	-> Gather Merge (cost=4068.62..4097.26 rows=249 width...	
17	Workers Planned: 1	
18	Workers Launched: 1	
19	-> Sort (cost=3068.61..3069.23 rows=249 width=...	
20	Sort Key: member_1.country	
21	Sort Method: quicksort Memory: 43kB	
22	Worker 0: Sort Method: quicksort Memory:...	
23	-> Partial HashAggregate (cost=3056.21...	
24	Group Key: member_1.country	
25	Batches: 1 Memory Usage: 93kB	
26	Worker 0: Batches: 1 Memory Usage:...	
27	-> Parallel Seq Scan on member memb...	
28	Planning Time: 0.321 ms	
29	Execution Time: 74.956 ms	

Resultados: FAV esa la mejor opcion por simplicidad y velocidad

ejemplo 3

Se desea conocer el titulo, fecha de creacion, formato, genero y promedio de duracion de las ultimas 5 canciones anadidas a su catalogo, y la duracion total segun el genero al que pertenezcan

SQL Basico

```
select title,
        created_at,
        format,
        genre,
        (select avg(last_five_documents.duration)
         from (select duration
                from document d2
                join media m2 on d2.id_document = m2.id_document
                where d2.created_at <= d1.created_at
                order by d2.created_at desc
                limit 5) as last_five_documents) as average_duration,
        (select sum(m2.duration)
```

```
        from document d2
            join media m2 on d2.id_document = m2.id_document
        where m2.genre = m1.genre)          as total_per_genre
from document d1
    join media m1 on d1.id_document = m1.id_document;
```

FAV

```
select title,
        created_at,
        format,
        genre,
        avg(duration) over (order by created_at desc rows between 5 preceding and
current row ) as average_duration,
        sum(duration) over (partition by genre)
as average_duration_per_genre
from document
    join media on document.id_document = media.id_document;
```



```
media.id_document),  
    total_duration_per_genre as (select genre, sum(duration) as total_duration  
                                from media  
                                group by genre)  
  
select title,  
       created_at,  
       format,  
       ad.genre,  
       average,  
       total_duration  
from document d  
    join average_duration ad on d.id_document = ad.id_document  
    join total_duration_per_genre td on ad.genre = td.genre;
```

	❏ QUERY PLAN	⌵
1	Hash Join (cost=34959.83..66921.59 rows=15725 width=198) (actual time=...	
2	Hash Cond: ((ad.genre)::text = (td.genre)::text)	
3	→ Hash Join (cost=33927.89..65848.36 rows=15725 width=190) (actual...	
4	Hash Cond: (d.id_document = ad.id_document)	
5	→ Seq Scan on document d (cost=0.00..31099.07 rows=177107 wi...	
6	→ Hash (cost=33731.33..33731.33 rows=15725 width=111) (actua...	
7	Buckets: 16384 Batches: 1 Memory Usage: 2034kB	
8	→ Subquery Scan on ad (cost=33298.89..33731.33 rows=15...	
9	→ WindowAgg (cost=33298.89..33574.08 rows=15725 ...	
10	→ Sort (cost=33298.89..33338.21 rows=15725...	
11	Sort Key: document.created_at DESC	
12	Sort Method: quicksort Memory: 2157kB	
13	→ Hash Join (cost=638.81..32202.80 r...	
14	Hash Cond: (document.id_document ...	
15	→ Seq Scan on document (cost=0...	
16	→ Hash (cost=442.25..442.25 ro...	
17	Buckets: 16384 Batches: 1 ...	
18	→ Seq Scan on media (cos...	
19	→ Hash (cost=835.38..835.38 rows=15725 width=83) (actual time=14.6...	
20	Buckets: 16384 Batches: 1 Memory Usage: 1909kB	
21	→ Subquery Scan on td (cost=520.88..835.38 rows=15725 width=...	
22	→ HashAggregate (cost=520.88..678.12 rows=15725 width=...	
23	Group Key: media_1.genre	
24	Batches: 1 Memory Usage: 3345kB	
25	→ Seq Scan on media media_1 (cost=0.00..442.25 r...	
26	Planning Time: 0.410 ms	
27	Execution Time: 235.391 ms	

Resultados: FAV esa la mejor opcion por simplicidad y velocidad

ejemplo 4

Se desea conocer los datos de profesionales que tienen deudas y el total de deuda relacionada a empleados que trabajan para la misma empresa

SQL basico

```
select f.id_fine,  
       f.id_document,  
       f.id_service,  
       p.organization,  
       (select sum(fee)  
        from fine f2  
         join loan_professional lp2 on f2.id_service = lp2.id_service  
         and f2.id_document = lp2.id_document  
         join professional p2 on lp2.id_member = p2.id_member  
         and f2.id_document = lp2.id_document  
        where p2.organization = p.organization) as total_fee  
from fine f  
     join loan_professional lp on f.id_service = lp.id_service
```

```
and f.id_document = lp.id_document
    join professional p on lp.id_member = p.id_member
and f.id_document = lp.id_document;
```

FAV

```
select id_fine,
       f.id_document,
       f.id_service,
       p.organization,
       sum(fee) over (partition by p.organization) as total_fee
from fine f
    join loan_professional lp on f.id_service = lp.id_service
and f.id_document = lp.id_document
    join professional p on lp.id_member = p.id_member
and f.id_document = lp.id_document;
```

	❏ QUERY PLAN	↕
1	WindowAgg (cost=4385.72..4385.99 rows=15 width=43) (actual time=280.3...	
2	→ Sort (cost=4385.72..4385.76 rows=15 width=39) (actual time=280...	
3	Sort Key: p.organization	
4	Sort Method: external merge Disk: 5280kB	
5	→ Nested Loop (cost=2985.41..4385.43 rows=15 width=39) (act...	
6	→ Hash Join (cost=2985.12..4377.85 rows=15 width=24) ...	
7	Hash Cond: ((lp.id_service = f.id_service) AND (lp...	
8	→ Seq Scan on loan_professional lp (cost=0.00...	
9	→ Hash (cost=1568.45..1568.45 rows=94445 width=...	
10	Buckets: 131072 Batches: 1 Memory Usage: 5...	
11	→ Seq Scan on fine f (cost=0.00..1568.45 ...	
12	→ Index Scan using professional_pkey on professional p...	
13	Index Cond: (id_member = lp.id_member)	
14	Planning Time: 0.407 ms	
15	Execution Time: 360.933 ms	

WITH

```

with total_fee_per_organization as (select organization, sum(fee) as total
                                   from fine f
                                   join loan_professional lp on
f.id_service = lp.id_service
                                   and f.id_document = lp.id_document
                                   join professional p on lp.id_member
                                   = p.id_member

```

```
                and f.id_document = lp.id_document
                group by organization)

select f.id_fine, f.id_document, f.id_service, tf.total
from fine f
    join loan_professional lp on f.id_service = lp.id_service
    and f.id_document = lp.id_document
    join professional p on lp.id_member = p.id_member
    and f.id_document = lp.id_document
    join total_fee_per_organization tf on p.organization = tf.organization;
```

QUERY PLAN

4	Hash Join (cost=5391.91..7668.71 rows=1 width=24) (actual time=320...
Select All	
	Cond: ((f.id_service = lp.id_service) AND (f.id_document = lp...
3	→ Seq Scan on fine f (cost=0.00..1568.45 rows=94445 width=16) (...)
4	→ Hash (cost=5391.83..5391.83 rows=5 width=16) (actual time=320...
5	Buckets: 16384 (originally 1024) Batches: 1 (originally 1) ...
6	→ Hash Join (cost=5226.40..5391.83 rows=5 width=16) (actu...
7	Hash Cond: (lp.id_member = p.id_member)
8	→ Seq Scan on loan_professional lp (cost=0.00..133...
9	→ Hash (cost=5226.10..5226.10 rows=24 width=12) (ac...
10	Buckets: 16384 (originally 1024) Batches: 1 (or...
11	→ Hash Join (cost=4386.32..5226.10 rows=24 wi...
12	Hash Cond: ((p.organization)::text = (tf.o...
13	→ Seq Scan on professional p (cost=0.00...
14	→ Hash (cost=4386.14..4386.14 rows=15 w...
15	Buckets: 8192 (originally 1024) Bat...
16	→ Subquery Scan on tf (cost=4385...
17	→ GroupAggregate (cost=4385...
18	Group Key: p_1.organizat...
19	→ Sort (cost=4385.72...
20	Sort Key: p_1.orga...
21	Sort Method: exter...
22	→ Nested Loop (...)
23	→ Hash Joi...
24	Hash C...
25	→ Se...
26	→ Ha...
27	...
28	...
29	→ Index Sc...

30	Index ...
31	Planning Time: 0.618 ms
32	Execution Time: 357.636 ms

Resultados: FAV y WITH estan casi empatados dado que la diferencia es demasiado pequena y las muestras fueron muy pocas

ejemplo 5

SQL Basico

```
select d.id_document,
       title,
       format,
       publication_place,
       (select avg(p2.dimension_height)
        from document d2
         join picture p2 on d2.id_document = p2.id_document
        where d2.publication_place = d.publication_place) as average_height,
       (select avg(p2.dimension_width)
        from document d2
         join picture p2 on d2.id_document = p2.id_document
        where d2.publication_place = d.publication_place) as average_width
```

```
from document d
    join picture p on d.id_document = p.id_document;
```

FAV

```
select document.id_document,
       title,
       format,
       publication_place,
       avg(dimension_width) over (partition by publication_place) as
average_width,
       avg(dimension_height) over (partition by publication_place) as
average_height
from document
    join picture on document.id_document = picture.id_document;
```

	QUERY PLAN	
1	WindowAgg (cost=33095.31..33409.09 rows=15689 width=159) (actual time...	
2	→ Sort (cost=33095.31..33134.54 rows=15689 width=103) (actual tim...	
3	Sort Key: document.publication_place	
4	Sort Method: quicksort Memory: 2429kB	
5	→ Hash Join (cost=438.00..32001.99 rows=15689 width=103) (a...	
6	Hash Cond: (document.id_document = picture.id_document)	
7	→ Seq Scan on document (cost=0.00..31099.07 rows=1771...	
8	→ Hash (cost=241.89..241.89 rows=15689 width=12) (act...	
9	Buckets: 16384 Batches: 1 Memory Usage: 803kB	
10	→ Seq Scan on picture (cost=0.00..241.89 rows=1...	
11	Planning Time: 0.235 ms	
12	Execution Time: 128.730 ms	

WITH

```

with averse_height as (select publication_place, avg(dimension_height) as
    avg_height
                        from document
                        join picture on document.id_document =
picture.id_document
                        group by publication_place),
    averse_width as (select publication_place, avg(dimension_width) as
    avg_width
                    from document

```



```
                                join picture on document.id_document =  
picture.id_document  
                                group by publication_place)  
select document.id_document, title, format, averate_width.publication_place,  
avg_height, avg_width  
from document  
    join averate_height on document.publication_place =  
averate_height.publication_place  
    join averate_width on document.publication_place =  
averate_width.publication_place;
```

	❏ QUERY PLAN
1	Hash Join (cost=65102.21..96815.48 rows=18251 width=159) (actual ti...
2	Hash Cond: ((document.publication_place)::text = (document_2.publi...
3	→ Hash Join (cost=32629.55..64193.57 rows=56854 width=139) (act...
4	Hash Cond: ((document.publication_place)::text = (average_he...
5	→ Seq Scan on document (cost=0.00..31099.07 rows=177107 w...
6	→ Hash (cost=32433.44..32433.44 rows=15689 width=44) (act...
7	Buckets: 16384 Batches: 1 Memory Usage: 750kB
8	→ Subquery Scan on average_height (cost=32080.43..3...
9	→ HashAggregate (cost=32080.43..32276.55 rows...
10	Group Key: document_1.publication_place
11	Batches: 1 Memory Usage: 2577kB
12	→ Hash Join (cost=438.00..32001.99 rows...
13	Hash Cond: (document_1.id_document = ...
14	→ Seq Scan on document document_1 ...
15	→ Hash (cost=241.89..241.89 rows=...
16	Buckets: 16384 Batches: 1 Me...
17	→ Seq Scan on picture (cost...
18	→ Hash (cost=32276.55..32276.55 rows=15689 width=44) (actual ti...
19	Buckets: 16384 Batches: 1 Memory Usage: 750kB
20	→ HashAggregate (cost=32080.43..32276.55 rows=15689 width...
21	Group Key: document_2.publication_place
22	Batches: 1 Memory Usage: 2577kB
23	→ Hash Join (cost=438.00..32001.99 rows=15689 width...
24	Hash Cond: (document_2.id_document = picture_1.i...
25	→ Seq Scan on document document_2 (cost=0.00...
26	→ Hash (cost=241.89..241.89 rows=15689 width=...
27	Buckets: 16384 Batches: 1 Memory Usage: ...
28	→ Seq Scan on picture picture_1 (cost=0...
29	Planning Time: 0.613 ms

Resultados: FAV esa la mejor opcion por simplicidad y velocidad

Investigacion sobre otras tecnicas de optimizacion no vistas en clase

Investigacion sobre otras tecnicas de optimizacion no vistas en clase (no puede ser fav, with, index)

1. identificarlas
2. caracterizadas
3. ejemplo

En esta seccion se hablara de algunas tecnicas de optimizaciones adicionales que son muy interesantes y unicas a su manera

- Parallel Queries
- Partitioning
- Materialized Views
- Prepared Statements

Otros son

- Prepared Statements
- indexes

Parallel Queries

Las consultas paralelas, o "parallel queries", son aquellas consultas SQL que se ejecutan de manera simultánea utilizando múltiples procesadores o núcleos del servidor de base de datos. Esta técnica de optimización se utiliza para acelerar el procesamiento de consultas complejas al dividir la carga de trabajo entre varios hilos de ejecución.

Es el trabajo del planificador de consultas determinar si es la mejor estrategia para usarse y cuantos nucleos puede usar, esto esta definido en la configuracion de postgres bajo la variable

`max_parallel_workers_per_gather`: El numero de workers que el planificador usara

Se caracteriza por:

- **División de Tareas:** Las tareas de la consulta se dividen en subconjuntos independientes que pueden ejecutarse simultáneamente.

- **Uso de Recursos Paralelos:** Se aprovechan múltiples núcleos o procesadores disponibles en el sistema para realizar operaciones en paralelo.
- **Coordinación de Resultados:** Al finalizar las operaciones paralelas, los resultados se combinan para producir el resultado final de la consulta.
- **Eficiencia en Consultas Complejas:** Se beneficia especialmente en consultas que involucran grandes conjuntos de datos o operaciones complejas que pueden dividirse en partes independientes.

Partitioning

El particionado es una técnica de optimización en bases de datos que consiste en dividir una tabla grande en secciones más pequeñas llamadas particiones. Cada partición tiene su propio conjunto de datos y, al dividir la tabla en partes más manejables, se pueden obtener beneficios en términos de rendimiento y mantenimiento. A continuación, se abordan los puntos clave sobre el particionado en SQL:

PostgreSQL tiene soporte para los siguientes 3 tipos de particiones:

- Range Partitioning
- List Partitioning
- Hash Partitioning

Tiene como características:

- **Mejora en el Rendimiento:** Las consultas y operaciones que se centran en un subconjunto específico de datos pueden ser más rápidas, ya que el sistema solo necesita buscar en la partición relevante en lugar de toda la tabla.
- **Facilita el Mantenimiento:** Las operaciones de mantenimiento, como la carga de datos y la eliminación, pueden ser más eficientes al trabajar con particiones más pequeñas en lugar de la tabla completa.
- **Optimización de Recursos:** Al limitar el acceso a ciertas particiones, el sistema puede asignar recursos de manera más efectiva para gestionar operaciones concurrentes.

```
CREATE TABLE mytable (id SERIAL PRIMARY KEY, created_at TIMESTAMP);  
CREATE TABLE mytable_2019 PARTITION OF mytable FOR VALUES FROM ('2019-01-01') TO  
('2020-01-01');  
CREATE TABLE mytable_2020 PARTITION OF mytable FOR VALUES FROM ('2020-01-01') TO  
('2021-01-01');  
CREATE TABLE mytable_2021 PARTITION OF mytable FOR VALUES FROM ('2021-01-01') TO  
('2022-01-01');
```

Materialized Views

Las vistas materializadas (Materialized Views) son objetos de base de datos que almacenan los resultados de una consulta para permitir un acceso más rápido y eficiente a los datos. A diferencia de las vistas regulares, que son consultas almacenadas y recalculadas cada vez que se llaman, las vistas materializadas almacenan físicamente los resultados y se actualizan periódicamente.

Se caracteriza por:

- **Almacenamiento Físico:** A diferencia de las vistas regulares, las vistas materializadas almacenan físicamente los datos en disco.
- **Actualización Programada:** Los datos en una vista materializada se actualizan periódicamente según una programación o en respuesta a eventos específicos.
- **Acceso Rápido:** Proporcionan un acceso más rápido a los resultados precalculados en comparación con la ejecución de la consulta original.

```
CREATE MATERIALIZED VIEW mv_ventas_mensuales AS
SELECT
    vendedor_id,
    EXTRACT(MONTH FROM fecha) AS mes,
    SUM(monto) AS total_ventas
FROM
```

```
ventas  
GROUP BY  
vendedor_id, EXTRACT(MONTH FROM fecha);
```

Bibliografia

- PostgreSQL Documentacion
 - [Parallel Querys](#)
 - [Partitioning](#)
 - [Materialized Views](#)
- Blog Article
 - [PostgreSQL Performance: Top 10 Optimization Tips | Medium](#)
- Youtube Video
 - [PostgreSQL Parallel Queries 12/12 - YouTube](#)

Conclusiones

Tras abordar distintas problemáticas del caso de estudio utilizando SQL básico, FAV y WITH, se ha observado que, en muchos casos, las window functions (FAV) ofrecen una solución más simple y eficiente. Esta técnica proporciona un rendimiento destacado al

cuando se usan funciones de agregacion en comparacion a no usarla. Además, se ha investigado sobre otras técnicas de optimización, como las consultas paralelas y el particionado de tablas, resaltando su impacto positivo en la mejora del rendimiento de las consultas en bases de datos.

La elección de la estrategia óptima depende de la naturaleza específica de la consulta y de la estructura de los datos. Las consultas paralelas resultan beneficiosas en operaciones complejas, mientras que el particionado de tablas es especialmente útil para mejorar la eficiencia en la gestión de grandes conjuntos de datos. En términos de simplicidad y velocidad, FAV se destaca en varias situaciones seguida por WITH que puede ser mas rapida en determinados casos, estas ofrecen una alternativa robusta para optimizar consultas en SQL.